

Your first week with EC2

(and other AWS things...)

AWS Summit 2013
Navigating the Cloud

Ryan Shuttleworth, Technical Evangelist



What are we going to cover?

Your first 5 days with EC2...

things you should know/think about up front ◀

some best practices for getting started ◀

essential technologies to dive into and get familiar with ◀

architectural principles you should immerse yourself in ◀

What are we going to cover?

Your first 5 days with EC2...

hear a 'looking back at our first year' customer story ◀
compressed into 5 days ◀

DAY

1

organise your
house



Users & Roles

Start as you mean to go on

Secure your console with IAM
roles

A little time spent now will save
headaches later

Users & Roles

Start as you mean to go on

Secure your console with IAM
roles

A little time spent now will save
headaches later

Accounts & Billing

Create a structure that makes sense

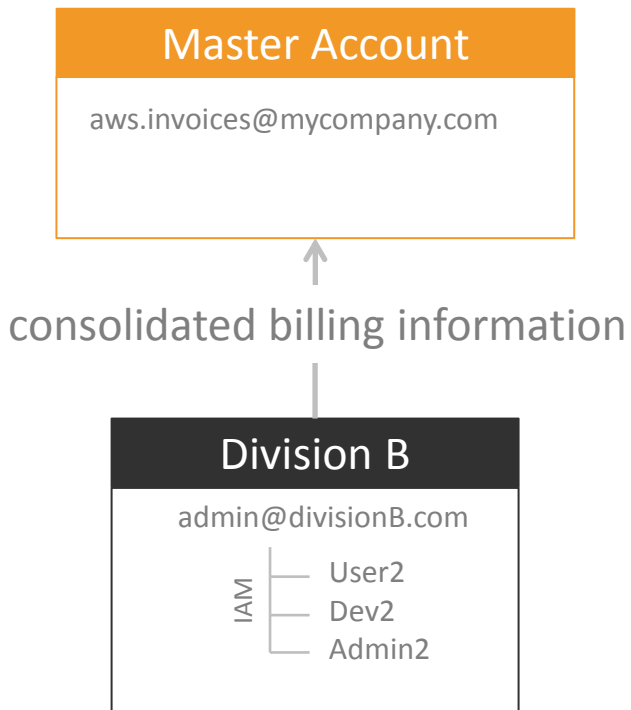
Dev & Test accounts vs production

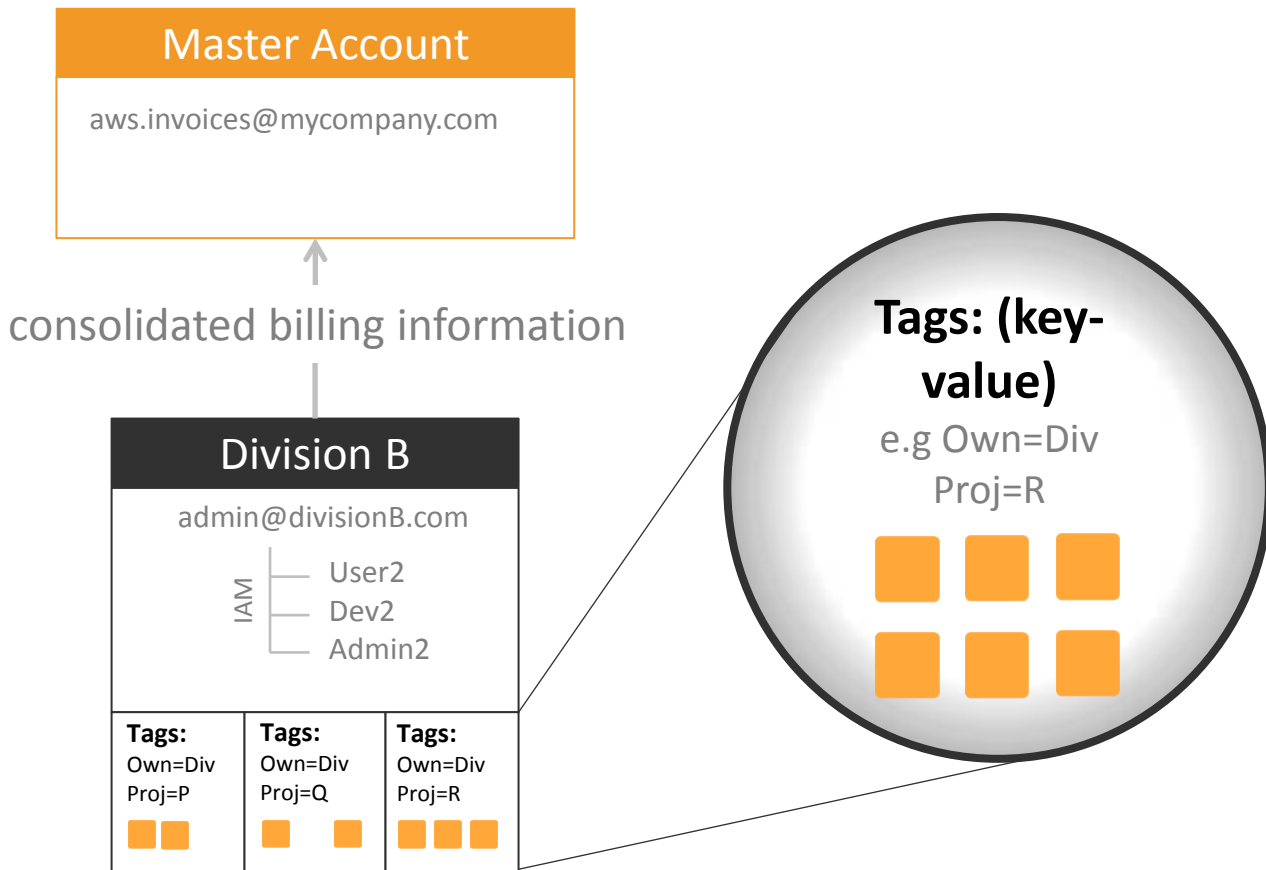
Consolidated billing

Resource tagging

Master Account

aws.invoices@mycompany.com





Master Account

aws.invoices@mycompany.com

consolidated billing information

Operating Co. A

admin@opcoa.com

IAM

- User1
- Dev1
- Admin1

Tags:Own=OpCo
Proj=A**Tags:**Own=OpCo
Proj=B**Tags:**Own=OpCo
Proj=C

Division B

admin@divisionB.com

IAM

- User2
- Dev2
- Admin2

Tags:Own=Div
Proj=P**Tags:**Own=Div
Proj=Q**Tags:**Own=Div
Proj=R

Business Unit C

admin@busUnitC.com

IAM

- User3
- Dev3
- Admin3

Tags:Own=BusC
Proj=X**Tags:**Own=BusC
Proj=Y**Tags:**Own=BusC
Proj=Z

Master Account

aws.invoices@mycompany.com

Operating

admin@opcoa.com

IAM

- User1
- Dev1
- Admin1

Tags:

Own=OpCo
Proj=A

Tags:

Own=OpCo
Proj=B

Tags:

Own=OpCo
Proj=C

Division

admin@divisionB.com


IAM

- User2
- Dev2
- Admin2

Tags:

Own=Div
Proj=P

Tags:

Own=Div
Proj=Q

Tags:

Own=Div
Proj=R

Business Unit

admin@busUnitC.com

IAM

- User3
- Dev3
- Admin3

Tags:

Own=BusC
Proj=X

Tags:

Own=BusC
Proj=Y

Tags:

Own=BusC
Proj=Z

Alert:

Reached \$500

Alert:

Reached \$3,000

Alert:

Reached \$1,250

Master Account

aws.invoices@mycompany.com

Programmatic billing access



S3

CSV



consolidated billing information

Operating Co. A

admin@opcoa.com

IAM

- User1
- Dev1
- Admin1

Tags:

Own=OpCo
Proj=A

Tags:

Own=OpCo
Proj=B

Tags:

Own=OpCo
Proj=C

Division B

admin@divisionB.com

IAM

- User2
- Dev2
- Admin2

Tags:

Own=Div
Proj=P

Tags:

Own=Div
Proj=Q

Tags:

Own=Div
Proj=R

Business Unit C

admin@busUnitC.com

IAM

- User3
- Dev3
- Admin3

Tags:

Own=BusC
Proj=X

Tags:

Own=BusC
Proj=Y

Tags:

Own=BusC
Proj=Z

Master Account

aws.invoices@mycompany.com

Programmatic billing access



S3

CSV



consolidated billing information

Operating Co. A

admin@opcoa.com

IAM

- User1
- Dev1
- Admin1

Tags:

Own=OpCo
Proj=A

Tags:

Own=OpCo
Proj=B

Tags:

Own=OpCo
Proj=C

Division B

admin@divisionB.com

IAM

- User2
- Dev2
- Admin2

Tags:

Own=Div
Proj=P

Tags:

Own=Div
Proj=Q

Tags:

Own=Div
Proj=R

Business Unit C

admin@bu3.com

IAM

- User3
- Dev3
- Admin3

Tags:

Own=BusC
Proj=X

Tags:

Own=BusC
Proj=Y

Cost	user:Owner	user:State
1.95	DbAdmin	Test
0.01	DbAdmin	Test
3.84	DbAdmin	Prod
6.00	DbAdmin	Test
234.63	SysEng	Prod
1.73	DbAdmin	Test
	DbAdmin	Prod
	DbAdmin	Prod

Secrets & Keys

Your front door keys



Secrets & Keys

Your front door keys

Secrets & Keys

Control access to
your instances

Key management
strategy

Your front door keys

Control access to
your APIs

Control access to
your instances

Secrets & Keys

Use IAM Roles to
distribute to instances

Key management
strategy

DAY **2**

learn the
basics



What is EC2?

Complete control

Flexible

Elastic capacity


Secure


Inexpensive


Reliable

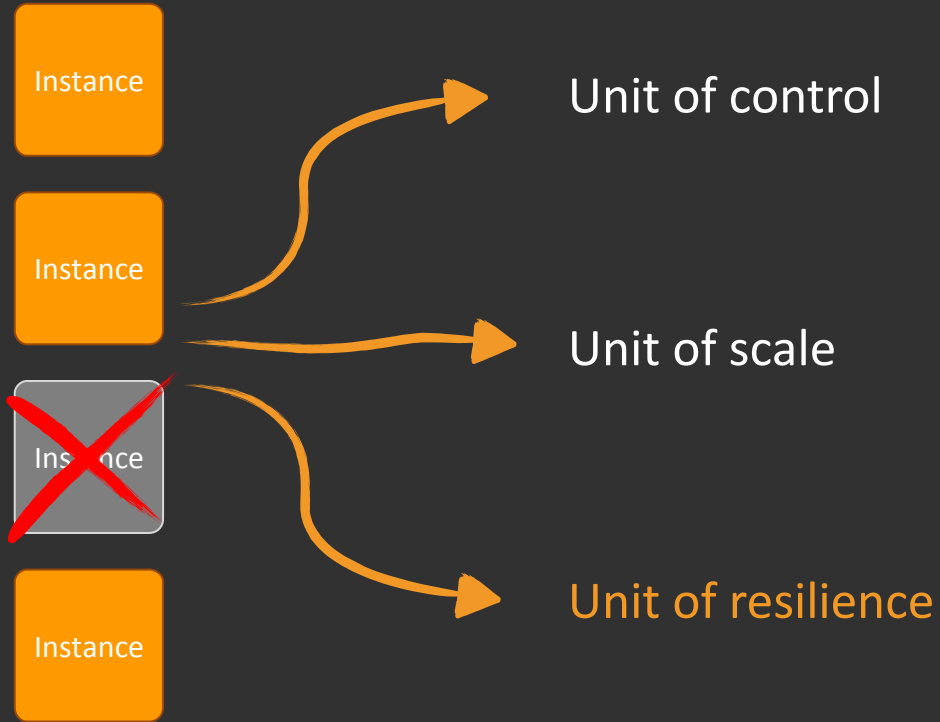
Disposable compute




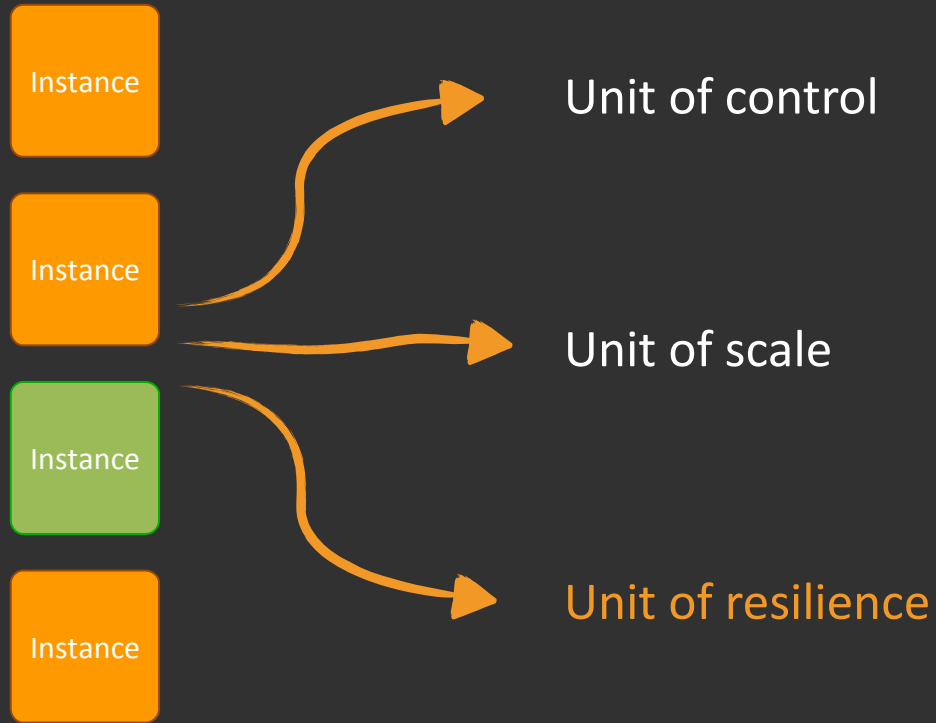












Think differently

Compute is transient

Programmatic
resources



Treat your datacentre
resources like code

Programmatic
resources



Treat your datacentre
resources like code

Distributed
systems



Design for decoupled
systems up front

Programmatic
resources



Treat your datacentre
resources like code

Distributed
systems



Design for decoupled
systems up front

Late binding



Decide what it will
run on at point of
deployment

DAY **3**

go wild with
tools

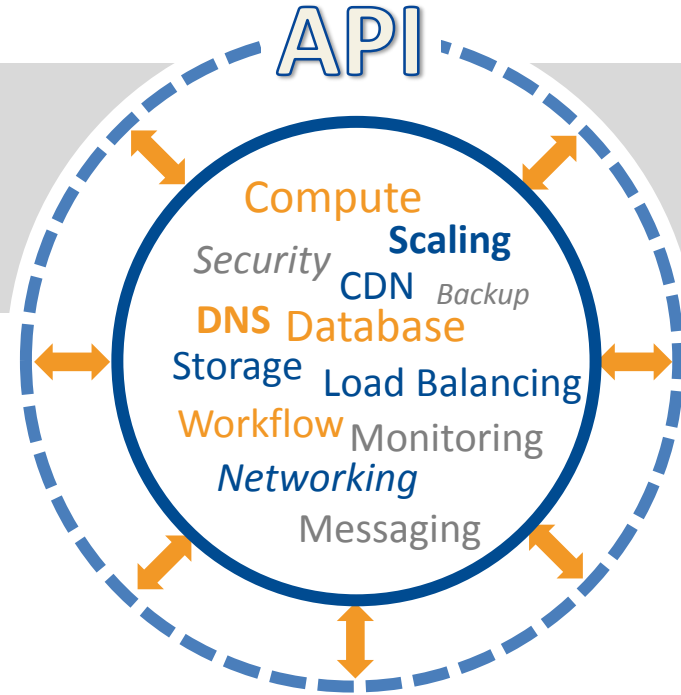


Everything is programmable

API

Access everything
via CLI, API or
Console

Achieve the highest levels
of automation
sophistication with ease



Microsoft
.net



```
$> ec2-run-instances ami-54cf5c3d  
--instance-count 2  
--group webservers  
--key mykey  
--instance-type m1.small
```

```
>>> import boto.ec2
>>> conn = boto.ec2.connect_to_region("us-east-1")
>>> conn.run_instances(
    'ami-54cf5c3d',
    key_name='mykey',
    instance_type='m1.small',
    security_groups=['webservers'])
```

Resources created programmatically

Resources created programmatically
Configure automatically

Bootstrapping

Bake an AMI

Start an instance

Configure the
instance

Create an AMI from
your instance

Start new ones from
the AMI

Bootstrapping

Bake an AMI

Start an instance

Configure the
instance

Create an AMI from
your instance

Start new ones from
the AMI

```
$> ec2-run-instances  
    <your ami-id>
```

Bootstrapping

Bake an AMI

vs

Configure dynamically

Start an instance

Configure the
instance

Create an AMI from
your instance

Start new ones from
the AMI

Launch an instance

Use metadata service
and cloud-init to
perform actions on
instance when it
launches

Bootstrapping

Bake an AMI

Build your base
images and setup
custom initialisation
scripts

Maintain your 'golden'
base



Configure dynamically

Use bootstrapping to
pass custom
information in and
perform post launch
tasks like pulling
code from SVN

Bootstrapping

Bake an AMI

Configure dynamically



Time consuming configuration
(e.g *startup time*)

Static configurations
(e.g *less change management*)

Bootstrapping

Bake an AMI

Configure dynamically



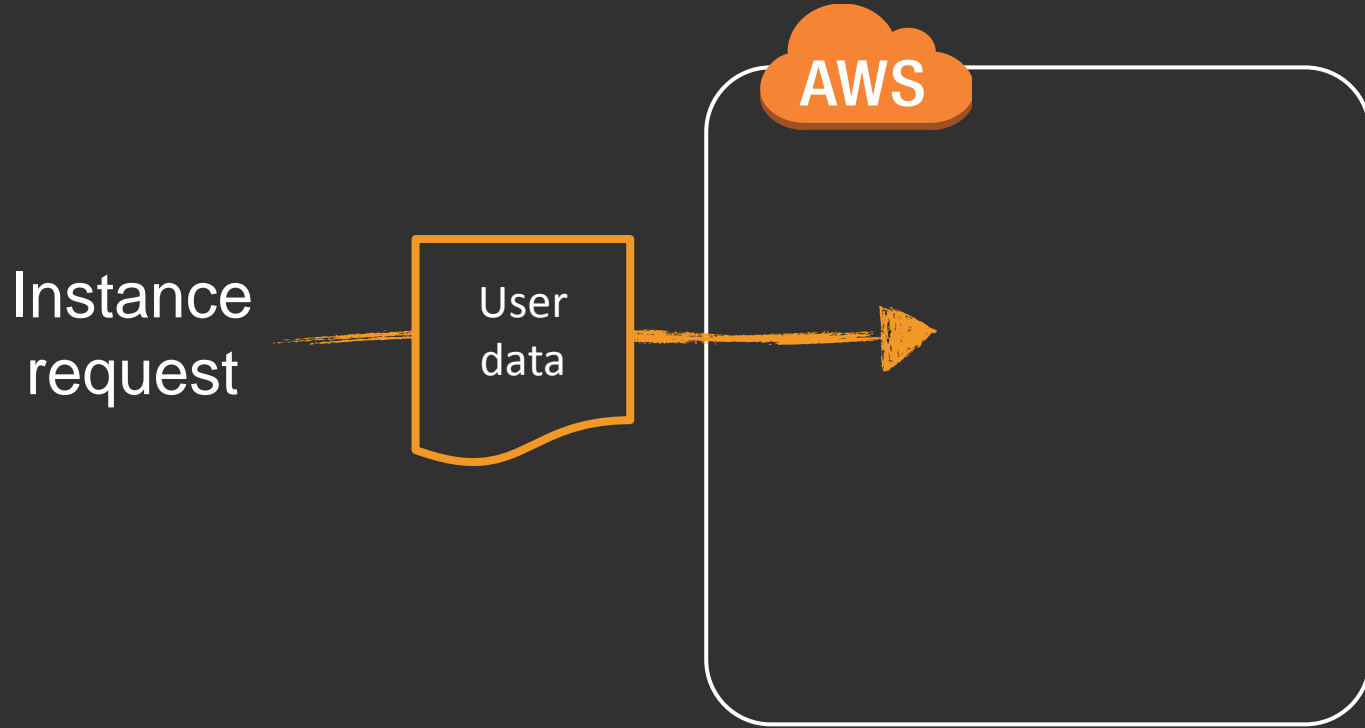
Continuous deployment
(e.g. *latest code*)

Environment specific
(e.g. *dev-test-prod*)

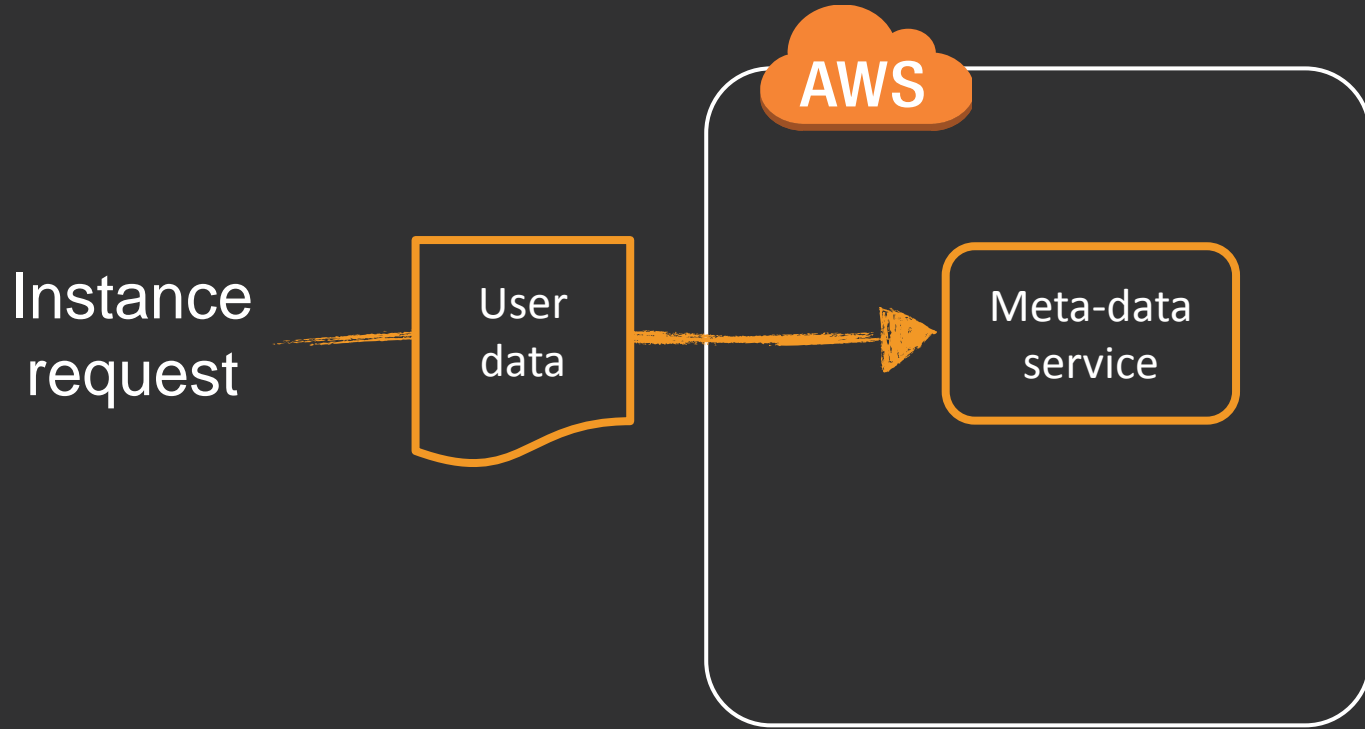
Goal is bring an instance up in a useful state

The balance will vary depending upon your application

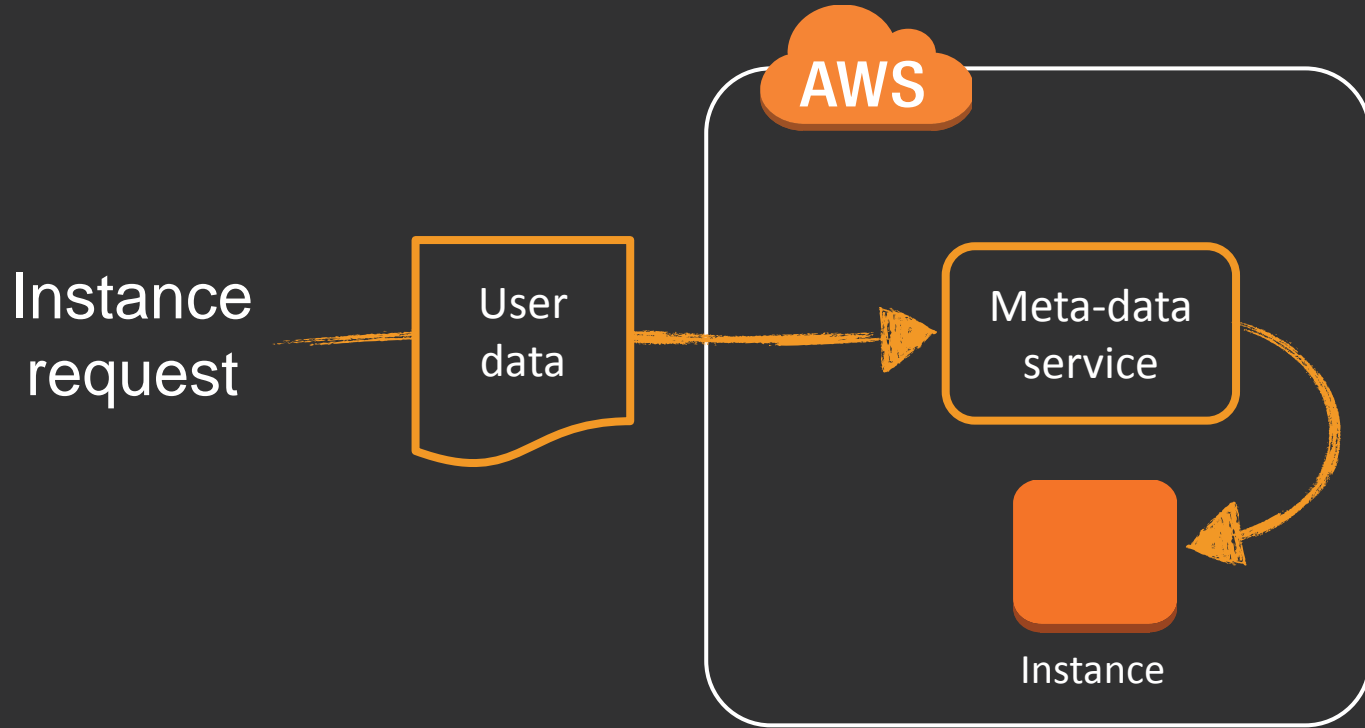
User-data



User-data



User-data



Shell script in user-data will be executed on launch:

```
#!/bin/sh
```

```
yum -y install httpd php mysql php-mysql  
chkconfig httpd on  
/etc/init.d/httpd start
```


Amazon Windows EC2Config Service executes user-data on launch:

```
<script>dir > c:\test.log</script>
```

```
<powershell>any command that you can run</powershell>
```

AWS Powershell Tools

```
<powershell>
```

```
    Read-S3Object -BucketName myS3Bucket  
    -Key myFolder/myFile.zip  
    -File c:\destinationFile.zip
```

```
</powershell>
```

Why do this?



Automation

Less fingers, less mistakes

Availability

Drive higher availability with self-healing

Efficiency

Audit and manage your estate with less time & effort

Scale

Manage large scale deployments and drive autoscaling

Flexible

Shell,
Powershell,
CloudFormation
, Chef, Puppet,
OpsWorks

Security

Instances
locked down by
default

DAY

4

try something
new



Unconstrained EC2 resources

Unconstrained Complimentary services




My little instance
(created programmatically)

A bit of S3 code

(pulled down automatically from S3)

```
>>> from boto.s3.key import Key
>>> k = Key(bucket)
>>> k.key = 'foobar'
>>> k.set_contents_from_string('This is a test of S3')
```




My little instance

(created programmatically)

A bit of S3 code

(installed automatically)


```
>>> from boto.s3.key import Key
>>> k = Key(bucket)
>>> k.key = 'foobar'
>>> k.set_contents_from_string('This is a test of S3')
```





Services instead of software

Removes undifferentiated heavy lifting



Services instead of software

Removes undifferentiated heavy lifting

S3 for object storage

SQS for queues

RDS for databases


CloudWatch for monitoring

DAY **5**

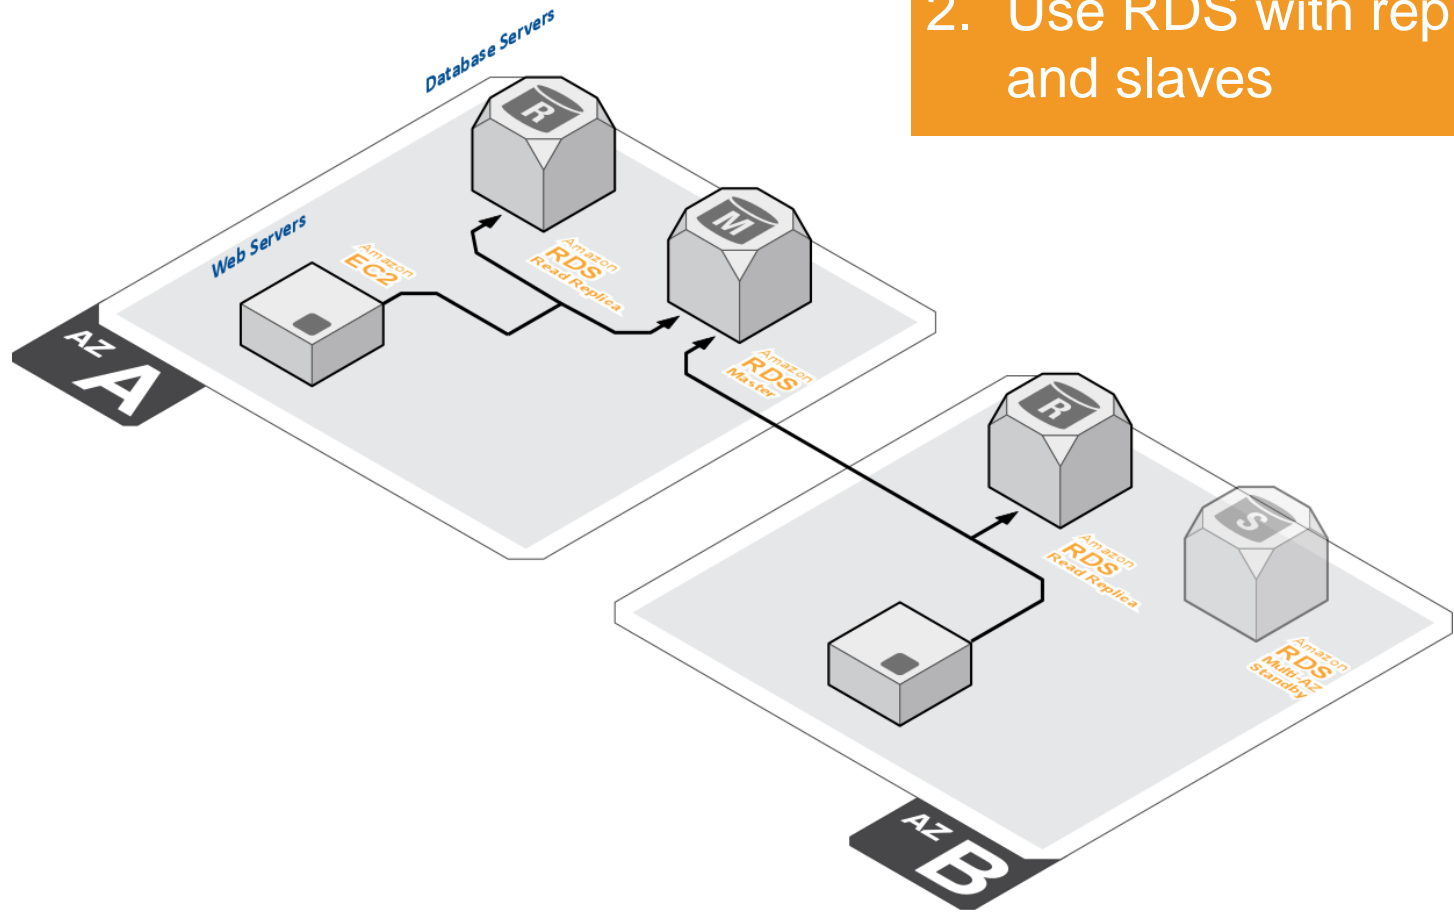
put something
together



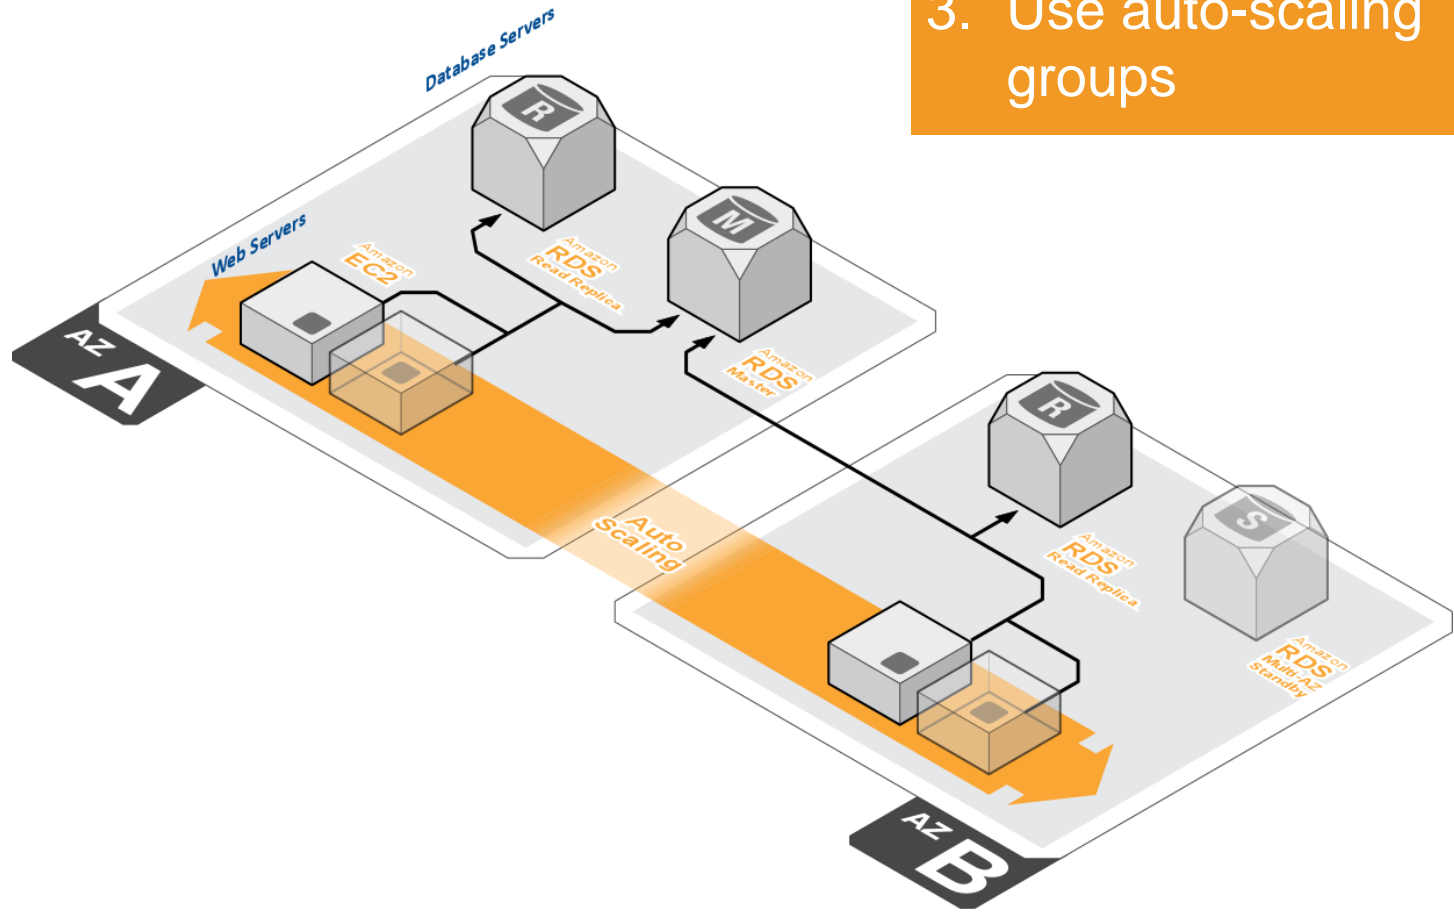
1. Use multiple availability zones



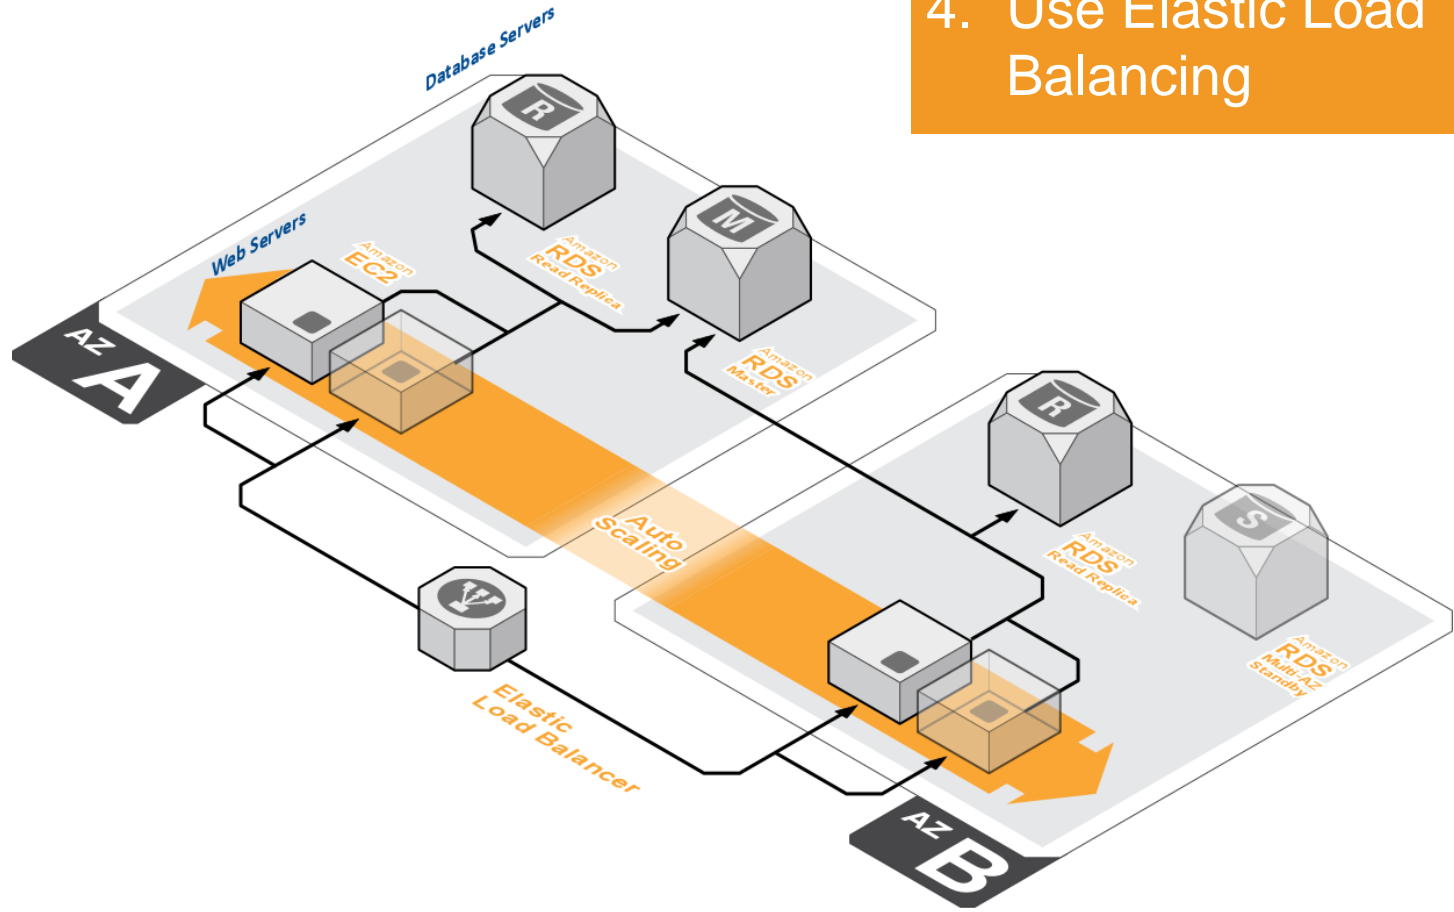
2. Use RDS with replicas and slaves



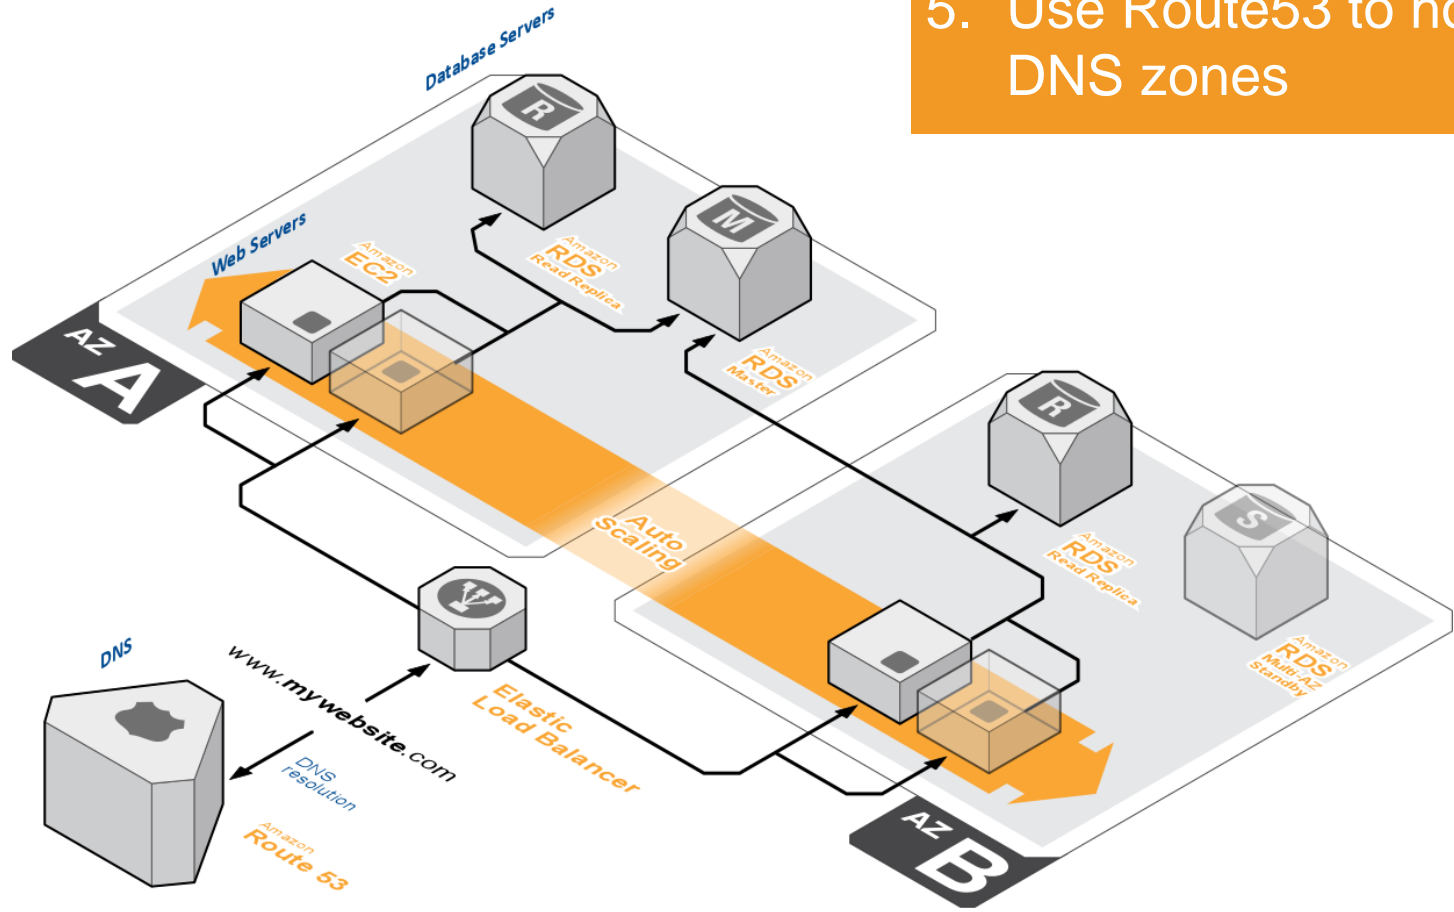
3. Use auto-scaling groups



4. Use Elastic Load Balancing



5. Use Route53 to host DNS zones



Find out more about model architectures in:

Building Web Scale Applications (bootstrapping track)

Architecting for High Availability (advanced track)



DAY

6

listen to Nick and Andy's story





Amazon Web Services

One Year, in 5 days...

What do we do



- Record peoples mobile and fixed communications
 - Voice, sms etc..
 - Highly sensitive data
 - Real time analytics and alerting
 - Over 120 Financial Services organisations and other verticals

What is Important



- Security
- Scalability
- Flexibility
- Global Platform

Day Zero




- Recently Funded
 - Full Pipeline of customers
 - 6 Weeks to 'go live'
-
- Web Console Service & Storage needed asap for demos

Day One – I found us a Server!



One “*slightly* fire damaged” server going cheap..

Day One – Getting Started




Day One – Lessons Learnt



- I'm dangerous on my own – I need a SysOps Guru
- Elastic IPs – Elastic Load Balancers & Route53
- RDS Monitoring – avoid that Sunday 'brown alert' moment, use cloud watch
- Single Instances... No
 - Stateless and multiple instances
- Decouple stuff, SQS, SWF, SNS

Day Two

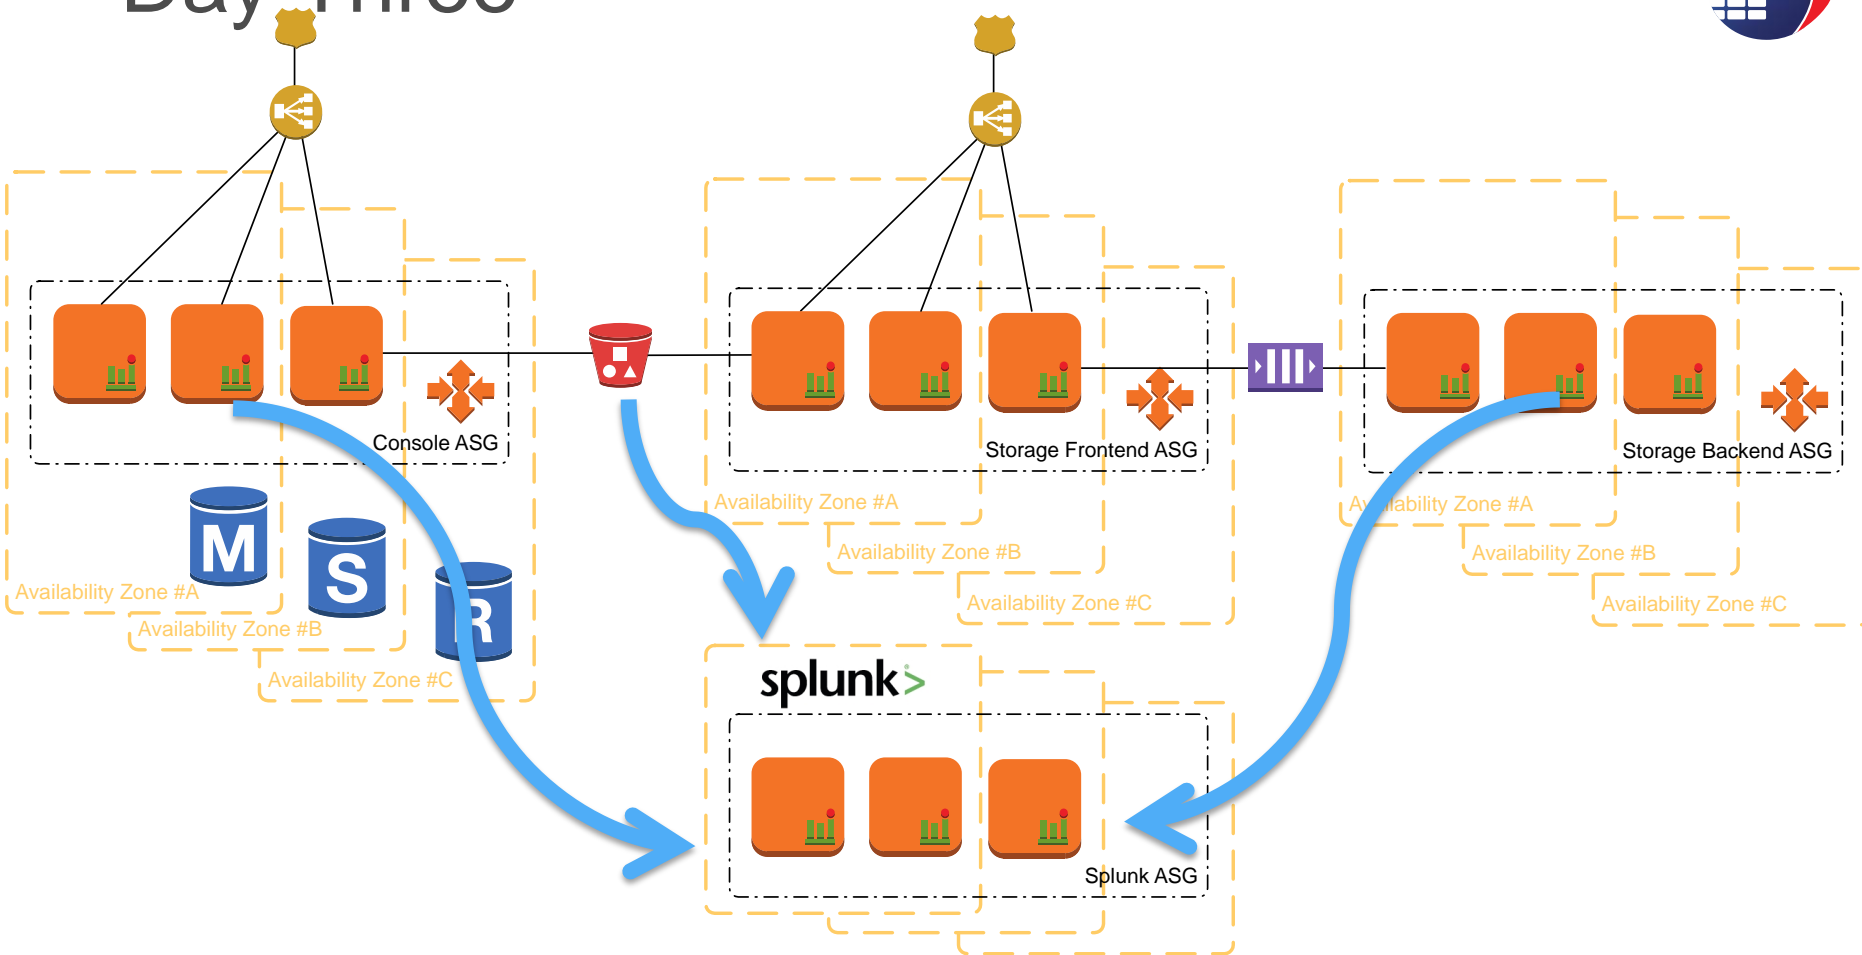


Day Two – Lessons Learnt



- Larger Estate:
 - Autoscaling
 - More Logs (evidence)
 - More Metrics
 - More code, more services, more features.
- Devolve everything (where possible) to be HTTP – easier to scale
- Make things stateless – accept failure as routine

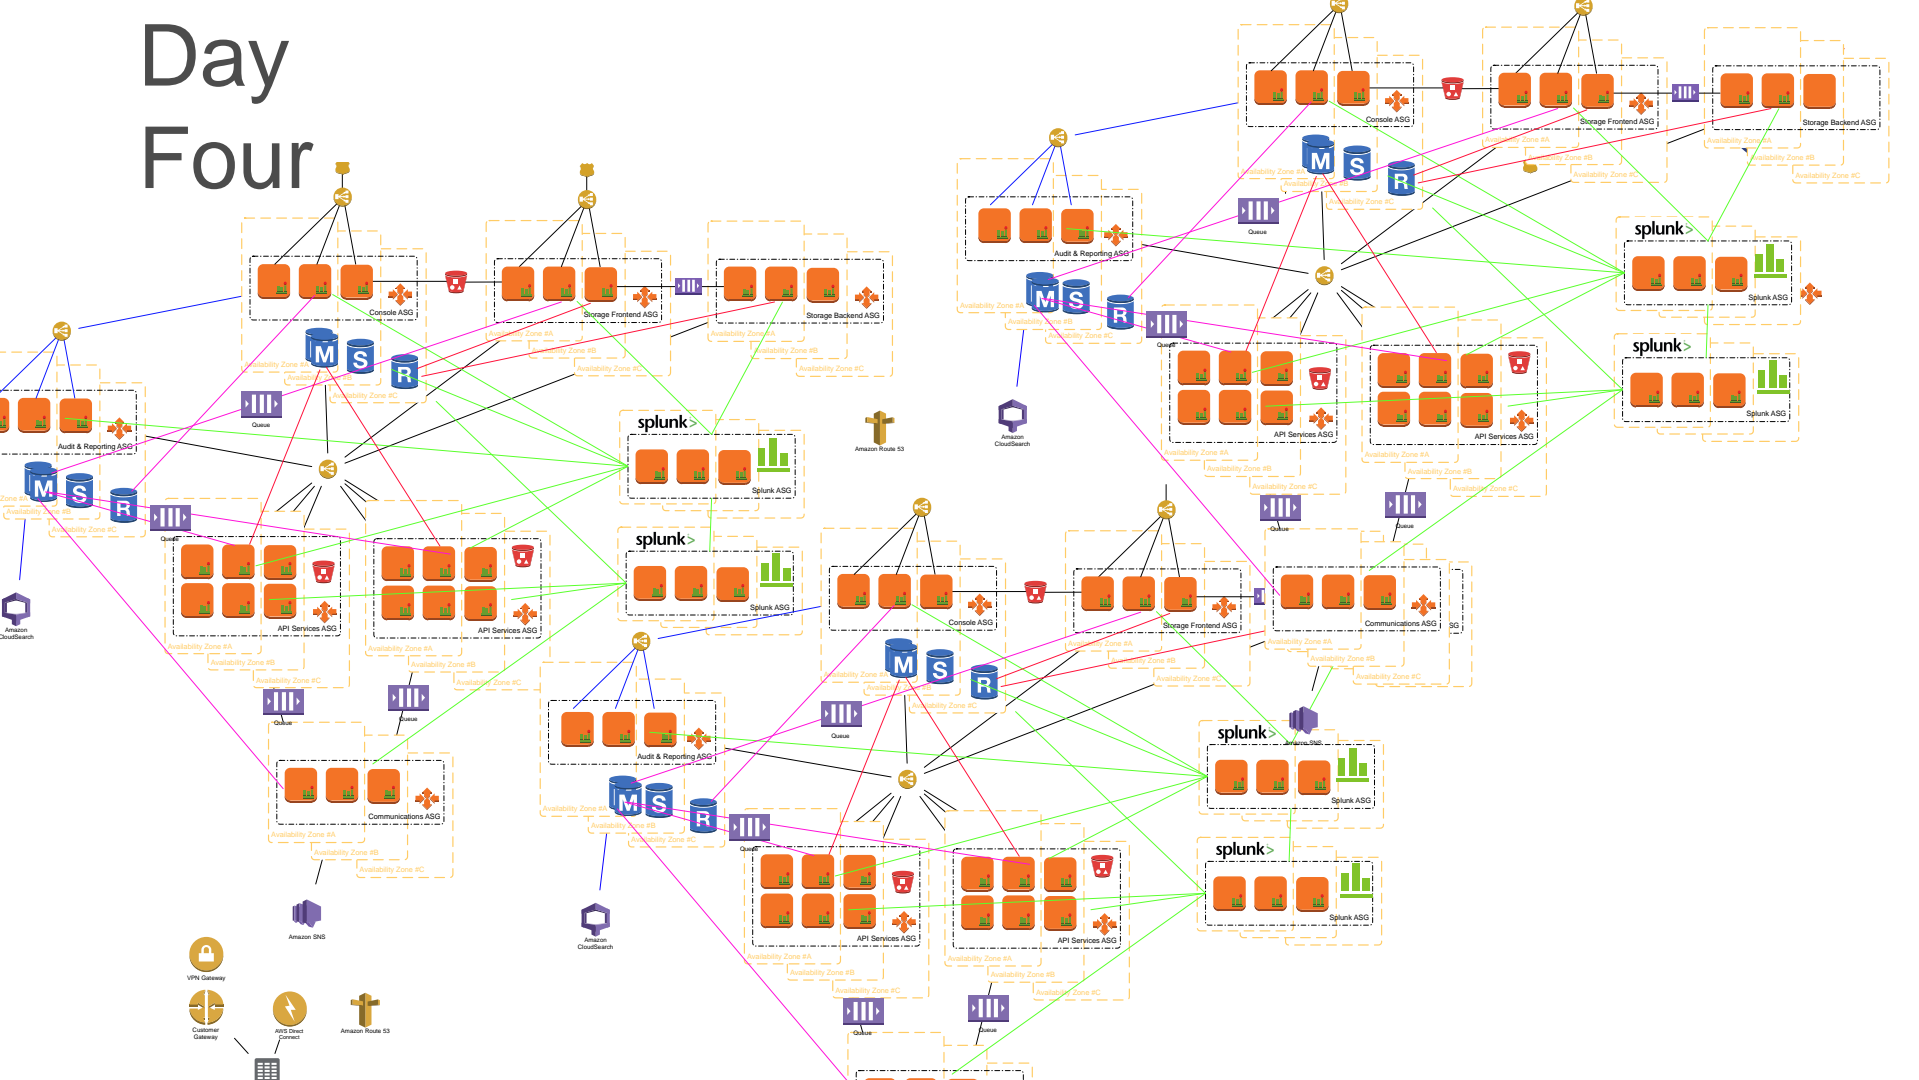
Day Three



Day Three – Lessons Learnt



- No human access to live
- Centralized Logging
 - We use Splunk, considered an employee in its own right
 - Consumes CloudWatch, S3 Logs, Application logs
 - Hand built alerts and filter by what is relevant
 - Evidence is key to diagnosis and resolution of any issues
- Amazon Support – cool story! Invaluable
- Think Big, you have a potential global platform at your fingertips.

[illegible]

Day Four– Lessons Learnt



- Check your growth – Easy to run more than you need.
- Document and diagram your system – rogue instances.
- Use Consolidated Billing
 - Separate accounts for Live & Test.
 - Get it off the credit card – Invoiced Billing
- Use as many of the tools as you can. - Don't re-invent the wheel, exploit the full ecosystem.
- Read the blogs, announcements, examples, and talk to the AWS SA's

Day Five



- Check out trusted advisor
 - What's your score?
- Leave work early!
 - You have a Global fault tolerant self healing system.
 - It notifies you if there is a problem, and then resolves itself.
 - You can keep an eye on your logs from the Pub.
- On a beer mat, do your reserved instance calculations
 - Buy reserved instances on Monday!

Lessons Learnt



- DevOps – tightly coupled development and systems teams = rapid evolution
- Building it right (evolution) allowed us to take time out and not be worried by failure
- Failure doesn't have to be all bad, if you expect it.
- Automation of testing, release and the deployment process, removes the risk from human mistake.



www.compliantphones.com

Summary

Lessons learned...

Monitoring, auto-scaling, de-coupling ◀

Stateless, expect failure ◀

Lock-down live, AWS support rocks, think big! ◀

Consolidated billing, follow the blog ◀

Reserved instances ◀

aws.typepad.com

